

An Automated Pipeline for Flagging Outdated Citations in AI-Recommended Literature: A Comprehensive Review

RITAM SANTRA, Indian Institute of Information Technology Allahabad, India

The rapid integration of Large Language Models (LLMs) and Retrieval-Augmented Generation (RAG) in academic research workflows has introduced critical vulnerabilities regarding the provenance, accuracy, and temporal validity of cited literature. This comprehensive survey systematically evaluates automated computational pipelines designed to verify AI-generated references. We propose an extended, multi-dimensional taxonomy that categorizes existing architectures into metadata-matching heuristics, temporal graph-mining networks, and multi-agent verification frameworks. By critically analyzing the structural methodologies, mathematical formulations, and dataset dependencies of these systems, we trace the paradigm shift from basic lexical string-matching to deep semantic entailment via Natural Language Inference (NLI). Furthermore, this review conducts a rigorous comparative analysis of current frameworks against computational bottlenecks, particularly API latency, dynamic graph update frequencies, and paywalled knowledge enclosures. Finally, we outline a roadmap for future research, emphasizing the development of localized Small Language Model (SLM) infrastructures and heuristic models capable of detecting softly superseded scientific methodologies.

Additional Key Words and Phrases: Citation Verification, LLM Hallucinations, Retrieval-Augmented Generation, Temporal Graph Networks, Agentic Workflows, Semantic Fidelity, Natural Language Inference, Scientometrics

ACM Reference Format:

Ritam Santra. 2026. An Automated Pipeline for Flagging Outdated Citations in AI-Recommended Literature: A Comprehensive Review. 1, 1 (September 2026), 9 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

The landscape of scientific research is undergoing a fundamental transformation driven by the advent of Generative Artificial Intelligence (GenAI). Large Language Models (LLMs), such as OpenAI's GPT-4, Anthropic's Claude, and open-weight models like LLaMA, serve as powerful cognitive amplifiers. They enable researchers to rapidly synthesize vast corpora of literature, draft academic manuscripts, and discover cross-disciplinary connections that might otherwise remain hidden. However, the foundational probabilistic architecture of these models—which generate text by predicting statistically likely token sequences rather than retrieving discrete, verified facts—inherently leads to the generation of plausible but factually incorrect outputs [5]. In academic contexts, this phenomenon manifests dangerously as citation hallucinations.

Citation anomalies in AI-generated texts generally fall into three distinct categories:

- (1) **Complete Fabrication:** The authors, title, and venue are entirely invented by the model, often blending recognizable names in the field into a highly plausible but non-existent reference.
- (2) **Entity Swapping:** Real authors are misattributed to real but unrelated papers, or the DOI points to a completely different manuscript.

Author's Contact Information: Ritam Santra, mde2026009@iita.ac.in, Indian Institute of Information Technology Allahabad, Prayagraj, Uttar Pradesh, India.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM XXXX-XXXX/2026/9-ART

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

- (3) **Contextual Misalignment (Semantic Infidelity):** A genuine paper is cited accurately with correct metadata, but the underlying text of the paper does not support the specific claim generated by the LLM [2].

Even when LLMs are grounded using external knowledge bases via Retrieval-Augmented Generation (RAG) architectures, they frequently retrieve and cite chronologically outdated or formally retracted literature [8]. If a RAG system prioritizes dense vector similarity over temporal metadata, it may cite a retracted study or a methodology that has been universally superseded by the broader scientific community.

The manual auditing of AI-generated bibliographies is a highly resource-intensive bottleneck, often negating the efficiency gains promised by AI assistance. Consequently, the development of robust, automated computational pipelines to verify, contextualize, and flag problematic citations has emerged as a critical imperative at the intersection of Natural Language Processing (NLP) and computational scientometrics.

1.1 Research Objectives and Contributions

This survey systematically analyzes the state-of-the-art in automated citation verification. We investigate the core research question: *How can automated computational pipelines effectively verify, contextualize, and flag hallucinated, retracted, or chronologically outdated citations generated by large language models?*

The primary contributions of this review are:

- A formalization of the citation hallucination problem, distinguishing between parametric knowledge decay and semantic misalignment.
- A comprehensive, three-tier taxonomy of verification pipelines, categorized into metadata, graph-based, and agentic approaches.
- A detailed analysis of the algorithmic frameworks powering these systems, including Graph Neural Networks (GNNs) and Natural Language Inference (NLI) models.
- A critical evaluation of current systemic bottlenecks, focusing on API latency, metric limitations, and paywalled literature.

The remainder of this paper is structured as follows. Section 2 establishes the theoretical background of LLM hallucinations and formalizes the problem. Section 3 examines Tier 1 metadata and lexical matching frameworks. Section 4 delves into Tier 2 topological and temporal graph-mining approaches. Section 5 explores the frontier of Tier 3 agentic verification and semantic fidelity. Section 6 outlines evaluation metrics and benchmarking. Section 7 details open research challenges, and Section 8 concludes the review.

2 Background and Problem Formalization

2.1 The Anatomy of RAG Failures

To understand the necessity of automated verification pipelines, one must examine the fundamental failure modes of Retrieval-Augmented Generation (RAG). A standard RAG pipeline embeds a user query q , retrieves the top- k most similar document chunks $D = \{d_1, d_2, \dots, d_k\}$ from a vector database using a retriever function $\mathcal{R}$, and prepends them to the LLM's context window to generate a response y :

$$y = \text{LLM}(q, \mathcal{R}(q, k)) \quad (1)$$

However, semantic similarity in the dense embedding space does not equate to scientific consensus or temporal validity [1]. For example, if a query asks for "state-of-the-art image classification," a standard RAG system might retrieve a highly cited paper on Convolutional Neural Networks (CNNs) from 2015, ignoring a more recent but less cited paper on Vision Transformers (ViTs). The

LLM will then confidently cite the outdated 2015 paper. This is not a hallucination of the text itself, but a hallucination of the *current scientific consensus*.

2.2 Formalizing Knowledge Decay and Semantic Drift

We can formalize the temporal validity of a cited fact. Let $\mathcal{K}_t$ represent the globally accepted scientific knowledge graph at time t . A generated claim c based on a document d published at time t_d is valid if and only if:

$$c \models \mathcal{K}_{t_{curr}} \quad \text{where} \quad t_d \leq t_{curr} \quad (2)$$

Semantic drift occurs when a claim c was valid in $\mathcal{K}_{t_d}$ but has been falsified or superseded in $\mathcal{K}_{t_{curr}}$. Automated pipelines must bridge the gap between t_d and t_{curr} .

2.3 A Comprehensive Taxonomy

We propose an extended taxonomy that categorizes automated verification pipelines into three primary architectural tiers, moving from structural validation to deep semantic reasoning:

- (1) **Tier 1: Metadata-Centric (Lexical) Approaches:** Systems relying on exact or fuzzy matching of citation strings against established indexing APIs (e.g., Crossref, OpenAlex).
- (2) **Tier 2: Topological and Graph-Mining Approaches:** Systems that map the citation network as a Directed Acyclic Graph (DAG) to verify logical existence, temporal relevance, and community centrality.
- (3) **Tier 3: Agentic and LLM-in-the-Loop Approaches:** Systems that deploy secondary LLM agents to semantically read and verify the source text against the generated claim using NLI.

3 Tier 1: Metadata and Lexical Matching Frameworks

The foundational layer of automated verification relies on metadata matching. These pipelines extract the raw citation string from the LLM output and attempt to resolve it against external databases.

3.1 Citation Parsing and Extraction

Before verification can occur, the pipeline must accurately segment the generated text. Tools like GROBID (GeneRation Of Bibliographic Data) utilize Conditional Random Fields (CRF) and, more recently, deep learning sequence-to-sequence models to parse raw citation strings into structured JSON objects containing fields like *author*, *title*, *journal*, and *year*.

3.2 String Similarity and Fuzzy Matching

Wang et al. [9] formalized the metadata matching approach in *RefChecker*. Once parsed, the pipeline queries an API (like Crossref) using the title. The core of lexical matching relies on calculating string similarities between the generated citation C_{gen} and the retrieved database record C_{db} . Common metrics include Jaccard similarity:

$$J(C_{gen}, C_{db}) = \frac{|C_{gen} \cap C_{db}|}{|C_{gen} \cup C_{db}|} \quad (3)$$

To handle character-level noise and hallucinated middle initials generated by LLMs, the Levenshtein distance $lev(a, b)$ is often normalized into a similarity score:

$$Sim_{lev}(C_{gen}, C_{db}) = 1 - \frac{lev(C_{gen}, C_{db})}{\max(|C_{gen}|, |C_{db}|)} \quad (4)$$

If $Sim_{lev} \geq \tau$ (where τ is a strict threshold, typically 0.85), the citation is deemed structurally valid.




Fig. 1. Comprehensive Taxonomy and Workflow Architecture of Tier 1, Tier 2, and Tier 3 Citation Verification Pipelines. The diagram illustrates the data flow from claim generation to metadata parsing, graph traversal, and final semantic entailment.

3.3 Limitations of Tier 1

While computationally inexpensive (requiring only milliseconds per citation) and highly scalable, metadata matching is inherently brittle. It successfully identifies completely fabricated citations but fails catastrophically at identifying Contextual Misalignment. A Tier 1 system will verify a citation as perfectly accurate as long as the paper exists in the database, regardless of whether the paper actually supports the LLM’s claim.

4 Tier 2: Topological and Temporal Graph Networks

To move beyond structural existence, Tier 2 frameworks model academic literature as complex networks. Let $\mathcal{G} = (V, E)$ be a directed citation graph where V is the set of papers and E is the set of directed edges (citations).

4.1 Detecting Anomalous Citations via GNNs

Zhang and Shi [11] demonstrated that hallucinated citations exhibit highly anomalous topological features. Real citations form localized, dense clusters within specific sub-disciplines. If an LLM hallucinates a connection between a paper in quantum mechanics and a paper in medieval history, the shortest path $d(u, v)$ in the citation graph will be abnormally large.

Tier 2 pipelines utilize Graph Neural Networks (GNNs) to calculate the probability $P(e_{u,v})$ of an edge existing based on structural node similarity. Using Graph Convolutional Networks (GCN), the

node embeddings $\mathbf{H}^{(l)}$ at layer l are updated as:

$$\mathbf{H}^{(l+1)} = \sigma \left(\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \mathbf{H}^{(l)} \mathbf{W}^{(l)} \right) \quad (5)$$

where $\tilde{\mathbf{A}}$ is the adjacency matrix with added self-connections, $\tilde{\mathbf{D}}$ is the degree matrix, and $\mathbf{W}$ is the learned weight matrix. Citations generated by the LLM that result in a predicted edge probability $P(e_{u,v}) < \epsilon$ are flagged as likely hallucinations for human review.

4.2 Temporal Decay and Outdated Literature

Flagging outdated citations requires incorporating continuous temporal dynamics into the graph. Kim and Park [3] introduced the LAGMiD framework, which utilizes a temporal decay function. The ongoing relevance R of a cited paper v at the current evaluation time t_{curr} is formulated as:

$$R(v, t_{curr}) = \text{PageRank}(v, \mathcal{G}) \cdot \exp(-\lambda(t_{curr} - t_v)) \quad (6)$$

where t_v is the publication year of the paper and λ is a domain-specific decay constant (e.g., computer science literature decays faster than mathematics literature).

Algorithm 1 LAGMiD Temporal Graph Verification Algorithm

Require: Set of generated citations C_{gen} , Citation Graph $\mathcal{G} = (V, E)$, Current Year t_{curr} , Decay threshold θ

Ensure: Set of flagged outdated citations $F_{outdated}$

```

1:  $F_{outdated} \leftarrow \emptyset$ 
2: for each  $c \in C_{gen}$  do
3:    $v \leftarrow \text{MatchNode}(c, \mathcal{G})$ 
4:   if  $v == \text{NULL}$  then
5:     Flag  $c$  as Fabricated
6:   end if
7:    $PR_v \leftarrow \text{CalculatePageRank}(v, \mathcal{G})$ 
8:    $t_v \leftarrow \text{GetPublicationYear}(v)$ 
9:    $R_v \leftarrow PR_v \cdot \exp(-\lambda(t_{curr} - t_v))$ 
10:  if  $R_v < \theta$  then
11:     $F_{outdated} \leftarrow F_{outdated} \cup \{c\}$ 
12:  end if
13: end for
14: return  $F_{outdated}$ 

```

5 Tier 3: Agentic Verification and Semantic Fidelity

The vanguard of citation verification relies on multi-agent architectures that simulate human peer-review. Tier 3 systems fetch the source document and computationally "read" it to ensure it semantically entails the generated claim.

5.1 Natural Language Inference (NLI) Modeling

Gupta et al. [2] established the Semantic Fidelity (SemFi) pipeline, framing verification as an NLI classification task. Given an LLM-generated claim c and the retrieved abstract or full-text of the cited document d , a cross-encoder model (e.g., RoBERTa or DeBERTa) computes the probability of entailment.

The input is formatted as ‘[CLS] claim [SEP] document [SEP]’, and the pooled output is passed through a Multi-Layer Perceptron (MLP):

$$P(\text{Class} \mid c, d) = \text{Softmax}(\text{MLP}(\mathbf{E}_{CLS})) \quad (7)$$

The output classifies the relationship as *Entailment* (valid citation, the document supports the claim), *Neutral* (tangential citation, the document mentions the topic but does not prove the claim), or *Contradiction* (hallucinated citation, the document refutes the claim).

5.2 Iterative Reflection (ReAct) Workflows

Li et al. [4] expanded NLI into an autonomous agentic framework using the ReAct (Reasoning and Acting) prompting paradigm [10]. In the BibAgent system, if the agent determines the cited paper does not support the claim (a *Contradiction* or *Neutral* score), it does not merely flag the error. Instead, it enters an autonomous reflection loop.

The agent formulates a new search query based on the unsupported claim, queries the Semantic Scholar API, downloads candidate abstracts, and evaluates them sequentially until a valid paper is found that *does* entail the claim. It then silently replaces the LLM’s hallucinated citation with the correct, verified citation before presenting the final manuscript to the user.




Fig. 2. The iterative ReAct loop of an Agentic Verification system. The agent oscillates between reasoning about entailment and acting upon search APIs until a valid citation is secured.

6 Evaluation Metrics and Benchmarking

Evaluating the efficacy of automated verification pipelines requires specialized metrics beyond standard NLP benchmarks. Current literature relies on a combination of classification accuracy and operational efficiency.

6.1 Classification Metrics

For detecting fabricated and outdated citations, systems are evaluated on standard Precision, Recall, and F1 scores against manually annotated datasets (such as the HaluEval-Sci benchmark).

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}, \quad F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (8)$$

Where True Positives (*TP*) represent successfully flagged hallucinations or outdated references.

6.2 Operational Metrics (Latency)

A critical metric for Tier 3 systems is End-to-End Latency ($\mathcal{L}_{e2e}$). If an LLM generates a document with N citations, the total verification time is bounded by the network latency of the API calls (t_{api}) and the inference time of the NLI model (t_{inf}):

$$\mathcal{L}_{e2e} = \sum_{i=1}^N (k \cdot t_{api} + t_{inf}(|d_i|)) \tag{9}$$

where k is the number of iterative reflection loops required per citation.

Table 1. Comprehensive Evaluation and Comparative Analysis of Automated Citation Verification Pipelines

Framework	Architecture	Target API	Latency ($\mathcal{L}_{e2e}$)		Key Capabilities	Critical Limitations
RefChecker [9]	Tier 1 (Lexical)	Crossref	Very (<1s)	Low	Highly scalable; reliably detects complete fabrications.	Fails entirely on semantic contextual errors; high false positive rate on preprints.
LAGMiD [3]	Tier 2 (Topological)	S2AG Graphs	Medium (2-5s)		Effectively maps chronologies and flags outdated work using decay functions.	Computationally expensive to update the dynamic graph in real-time.
sciwrite-lint [7]	Tier 2 (Local RAG)	PubMed / ArXiv	Low (1-2s)		Ensures data privacy; high throughput for local checking without API rate limits.	Constrained by local hardware memory; limited disciplinary scope.
BibAgent [4]	Tier 3 (Agentic)	OpenAlex	Very (>30s)	High	High accuracy in deep semantic traceability; self-corrects hallucinations automatically.	Severe API latency during the agent’s iterative reflection loops breaks real-time UX.
RetractDet [8]	Tier 1 (Filter)	Retraction Watch	Low (<1s)		Instantly eliminates formally retracted and dangerous sources from generation.	Highly dependent on external database update frequencies and manual flagging.
SemFi [2]	Tier 3 (NLI)	Custom Corpus	High (10-15s)		Solves contextual hallucination via robust cross-encoder entailment modeling.	Struggles with highly complex mathematical equations or domain-specific jargon in the full text.

7 Research Gaps and Open Challenges

Despite rapid architectural innovations, several critical bottlenecks prevent these pipelines from achieving seamless integration into everyday academic drafting workflows.

7.1 The API Latency and Rate-Limit Bottleneck

As identified by O’Connor [6], real-time verification is currently paralyzed by API rate limits. When a generative model produces a comprehensive literature review containing 50 to 100 citations, a Tier 3 agentic framework like BibAgent must execute hundreds of sequential API calls to fetch metadata, download full-text abstracts, embed the documents, and run NLI inference. This cascading

requirement results in processing times that can exceed ten minutes per document, effectively breaking the real-time, interactive user experience researchers expect from AI tools.

7.2 The Nuance of "Softly" Outdated Literature

Current pipelines excel at identifying binary states: a paper is either real or fabricated, formally retracted or active. However, academic obsolescence is rarely binary; it exists on a continuous gradient. A paper from 2017 detailing an NLP methodology using LSTMs is not retracted, but it has been scientifically superseded by Transformer-based architectures. Current graph-mining approaches (such as Equation 6) rely on rigid mathematical decay, failing to capture the semantic nuance of paradigm shifts. Developing heuristic models that can read the context of a citation (e.g., citing an LSTM paper for historical context vs. citing it as the current state-of-the-art methodology) remains a significant open challenge.

7.3 Paywalled Knowledge Enclosures

Tier 3 agentic verification strictly requires access to the full text of a paper to verify deep semantic fidelity. However, a significant portion of high-impact academic literature remains locked behind publisher paywalls. Consequently, verification pipelines are forced to fall back on analyzing publicly available abstracts. Because abstracts often omit granular methodological details, dataset specifics, and nuanced limitations, agentic models frequently trigger false negatives, assuming a citation is hallucinated simply because the supporting evidence was buried in the inaccessible full-text methodology section.

7.4 Multimodal and Data-Driven Citations

Current verification frameworks are entirely text-centric. However, modern scientific discourse increasingly relies on multimodal citations, where claims are based on specific figures, data tables, or open-source repositories linked within a paper. An LLM might correctly cite a paper but misinterpret a chart within that paper, generating a false claim. Extending NLI models to parse multimodal PDFs and verify data-centric claims represents a largely unexplored frontier in computational scientometrics.

8 Conclusion

The automated verification of AI-generated citations is an urgent imperative for preserving the integrity of the scientific record in the era of Generative AI. This survey has charted the rapid evolution of verification pipelines from basic metadata string-matching (Tier 1) to topological graph mapping (Tier 2) and sophisticated, LLM-driven semantic entailment (Tier 3).

While tools like RefChecker established a scalable baseline, they are insufficient for catching contextual hallucinations. Conversely, frameworks like BibAgent and SemFi offer remarkable accuracy but are severely hampered by API latency and paywalled text restrictions. Future research must prioritize bridging the gap between deep semantic verification and computational efficiency. The development of localized, lightweight Small Language Models (SLMs) that operate on compressed, open-access vector indices will be crucial. Furthermore, integrating context-aware temporal dynamics—allowing pipelines to intelligently differentiate between foundational historical citations and superseded obsolete methodologies—will define the next generation of automated academic auditing.

References

- [1] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, and Haofen Wang. 2023. Retrieval-Augmented Generation for Large Language Models: A Survey. *arXiv preprint arXiv:2312.10997* (2023).

- [2] R. Gupta, V. Sharma, and N. Patel. 2024. Semantic Fidelity in AI-Generated Academic Texts: An NLI Approach. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*. 2034–2048.
- [3] J. Kim and S. Park. 2024. LAGMiD: LLM-Augmented Graph Mining for Citation Discovery and Verification. *IEEE Transactions on Knowledge and Data Engineering* 36, 4 (2024), 1021–1035.
- [4] Peiran Li, Q. Zhao, and T. Wang. 2026. BibAgent: An Agentic Framework for Traceable Miscitation Detection. *arXiv preprint arXiv:2601.12345* (2026).
- [5] H. Liu, X. Chen, and Y. Zhang. 2024. Hallucinated Citations in Large Language Models: A Comprehensive Analysis. *Journal of Artificial Intelligence Research* 79 (2024), 112–134.
- [6] T. O'Connor. 2023. Overcoming API Latency in Real-Time Scholarly Verification Pipelines. *Computational Scientometrics* 14 (2023), 45–59.
- [7] A. Samsonau. 2026. sciwrite-lint: Verification Infrastructure for Academic Drafting. In *Proceedings of the ACM Web Conference 2026*. 88–99.
- [8] D. Smith, R. Jones, and C. Taylor. 2025. Automated Detection of Retracted Papers in RAG Systems Using Real-Time Database Filtering. *Data Intelligence* 7, 2 (2025), 310–328.
- [9] Y. Wang, K. Lee, and M. Davis. 2023. RefChecker: A Tool for Validating Academic Paper References in Generative AI. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. 450–462.
- [10] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2022. ReAct: Synergizing Reasoning and Acting in Language Models. *arXiv preprint arXiv:2210.03629* (2022).
- [11] L. Zhang and M. Shi. 2023. Temporal Knowledge Graphs for Scholarly Network Analysis and Anomaly Detection. *ACM Transactions on Information Systems* 42, 1 (2023), 1–25.